

**Model-based design of
at-the-edge heterogeneous Robotics Applications
based on ROS**

Candidate:
Simone Girardi

Thesis advisor:
Prof. Nicola Bombieri
Research supervisor:
Dott. Stefano Aldegheri

Abstract

Developing distributed robotics applications on embedded devices, we have to deal with the diversity of the applications and the different platforms where these applications run. At the state of the art there are some solutions that allow us to develop robotics applications and deploy them on embedded boards. The problem is that none of these solutions allows us to be sufficiently accurate to guarantee the functioning of the entire system, especially if we want to increase its complexity. To solve the problem we must take account of the necessary resources to run the applications and the constraints imposed by the limits of the devices.

Contents

Abstract	i
1 Introduction	1
1.1 Thesis outline	1
2 Background	3
2.1 Deep Learning	3
2.1.1 Motivations	3
2.1.2 Definitions	3
2.1.3 Performance Measurement	5
2.1.4 Frameworks	6
2.1.5 Challenges	7
2.2 Edge Computing	8
2.2.1 Motivations	8
2.2.2 Definitions	9
2.2.3 Performance Metrics	10
2.2.4 Challenges	10
2.3 Robotics Tools and Platforms	11
2.3.1 ROS	11
2.3.2 NVIDIA Isaac SDK	13
2.3.3 CoppeliaSim	13
2.4 Docker	14
2.4.1 Understanding Docker Concepts	14
2.4.2 Portability Limitations of Container Images	16
2.5 Kubernetes	17
2.5.1 What is Kubernetes	17
2.5.2 Architecture of a Kubernetes cluster	18
2.5.3 Running an application in Kubernetes	19

2.6	Kubeedge	21
3	Verification	23
3.1	System description	23
3.2	Unexpected behaviour	25
3.2.1	The problem of the non-determinism	26
3.2.2	The problem of the concurrency	26
3.3	Discussion	26
4	Methodology	27
4.1	A ROS based robotic system	27
4.2	First level architecture	27
4.3	Second level architecture	29
4.4	Third level architecture	30
5	Experimental Results	31
5.1	Setup	31
5.2	Mobile robots	32
5.2.1	Architecture at L1	32
5.2.2	Architecture at L2	33
5.2.3	Architecture at L3	33
5.3	Deployment from the cloud to the edge	33
5.3.1	VOVD at the edge	35
5.3.2	ORB-SLAM2 at the edge	35
5.4	The whole system on Kubeedge	35
5.5	Discussion	35
5.5.1	Hybrid local/global planner	35
5.6	Results	35
	Conclusions	39
	Bibliography	41

List of Figures

2.1	DNN example	4
2.2	Cloud computing paradigm	9
2.3	Edge computing paradigm	10
2.4	Example of nodes and topic communication	12
2.5	Docker images, registries, and containers	16
2.6	Kubernetes exposes the whole datacenter as a single deployment platform	18
2.7	The components that make up a Kubernetes cluster	19
2.8	A basic overview of the Kubernetes architecture and an application running on top of it.	20
3.1	Main blocks of the ORB-SLAM2 algorithm.	24
3.2	DAG of the feature extraction block and the corresponding sub-block implementations (GPU vs. CPU).	25
3.3	LOST states of ORB-SLAM2	26
4.1	L1 architecture	28
5.1	L1 architecture	32
5.2	L2 architecture	34
5.3	L3 architecture	34

List of Tables

2.1	Neural Network Perfomance Metrics	6
5.1	Technical specifications of the Host	31
5.2	Performance Measurement	36
5.3	Performance Measurement.	37

Chapter 1

Introduction

1.1 Thesis outline

This thesis is organised in two main workflows: *Verification* and *Methodology*. Chapter 2 is an overview over the research contribution in the field of deep learning, edge computing and the other tools discussed in the next chapter of this thesis. The first workflow is described in chapter 3, where a complete inspection and verification was made on the ORB SLAM algorithm to guarantee a deterministic behaviour in a sequential context. The second workflow described in chapter 4 proposes a solution to the problem of integrating heterogeneous robotic applications. With the support of edge computing platforms like Kubeedge, some experiments were made to deploy automatically our integrated system from the host (cloud) to the edge. Finally in chapter 5 the experimental results are shown and discussed.

Chapter 2

Background

2.1 Deep Learning

2.1.1 Motivations

Deep learning has recently been highly successful in machine learning across a variety of application domains, including computer vision, natural language processing, and big data analysis, among others. For example, deep learning methods have consistently outperformed traditional methods for object recognition and detection in the ISLVR Computer Vision Competition since 2012 [27]. However, deep learning’s high accuracy comes at the expense of high computational and memory requirements for both the training and inference phases of deep learning. Training a deep learning model is space and computationally expensive due to millions of parameters that need to be iteratively refined over multiple time epochs. Inference is computationally expensive due to the potentially high dimensionality of the input data (e.g., a high-resolution image) and millions of computations that need to be performed on the input data.

2.1.2 Definitions

As described in [12], the modern term ”deep learning” goes beyond the neuroscientific perspective engineering applications on the current breed of machine learning models. It appeals to a more general principle of learning *multiple levels of composition*, which can be applied in machine learning frameworks that are not necessarily neurally inspired. Deep learning is a subset of AI and machine learning and differs in that they can automatically




Figure 2.1: DNN example with image classification

learn representations from data such as images, video or text, to be used for classification without introducing hand-coded rules or human domain knowledge. Their highly flexible architectures can learn directly from raw data and can increase their predictive accuracy when provided with more data. A deep learning prediction algorithm, consists of a number of layers, as shown in Fig. 2.1.

In deep learning *inference*, the input data pass through the node's layers in sequence, and each layer performs matrix multiplications on the data. The output of a layer is usually the input to the subsequent layer. After data are processed by the final (fully connected) layer, the output is either a feature or a classification value. When the model contains many layers in sequence, the neural network is known as a deep neural network (DNN). When the matrix multiplications include convolutional filter operations, the model is named convolutional neural networks (CNNs), which is common for image and video processing contexts. There are also DNNs designed especially for time series prediction; these are called recurrent neural networks (RNNs), which have loops in their layer connections to keep state and enable predictions on sequential inputs.

In deep learning *training*, the computation proceeds in reverse order. Given the ground-truth training labels, multiple passes are made over the layers to optimize the parameters of each layer of matrix multiplications,

starting from the final layer and ending with the first layer. The algorithm used is typically stochastic gradient descent (SGD). In each pass, a randomly small subset of N input data ("mini-batch") from the training data set, is selected and used to update the gradients in the direction that minimizes the training loss (where the training loss is defined as the difference between the predictions and the ground truth). One pass through the entire training data set is called a training epoch [26].

There are some considerations to take into account: the first is that there are a large number of parameters in the matrix multiplications, resulting in many computations being performed and thus the latency issues that we see on end devices. The second is that there are many choices (hyper-parameters) on how to design the DNN models (e.g., the number of parameters per layer, and the number of layers), which makes the model design more of an art than a science. Different DNN design decisions result in tradeoffs between system metrics; for example, a DNN with higher accuracy likely requires more memory to store all the model parameters and will have higher latency because of all the matrix multiplications being performed. On the other hand, a DNN model with fewer parameters will likely execute more quickly and use less computational resources and energy, but it may not have sufficient accuracy to meet the application's requirements.

2.1.3 Performance Measurement

How can we evaluate the performance of a neural network? Deep learning can be used to perform both supervised learning and unsupervised learning. The metrics of success depend on the particular application domain where deep learning is being applied. For example, in object detection, the accuracy may be measured by the mean average precision (mAP) [27], which measures how well the predicted object location overlaps with the ground-truth location, averaged across multiple categories of objects. In machine translation, the accuracy can be measured by the bilingual evaluation understudy score metric [19], which compares a candidate translation with several groundtruth reference translations. Other general system performance metrics not specific to the application include throughput, latency, and energy. These metrics are summarized in Table 2.1. Designing a good DNN model or selecting the right DNN model for a given application is challenging due to the large number of hyperparameter decisions.

Metric	Unit
Latency	s
Energy	mW, J
Concurrent Requests Served	#
Network Bandwidth	Mbps
Accuracy	Application Specific

Table 2.1: Neural Network Performance Metrics

Machine learning research typically focuses on accuracy metrics, and their system performance results are often reported from powerful server testbeds equipped with GPUs. For example, Huang et al. [15] compared the speed and accuracy tradeoffs when running on a high-end gaming GPU (NVIDIA Titan X). The YOLO DNN model [24], which is designed for real-time performance, provides timing measurements on the same server GPU. Specifically targeting mobile devices, Lu et al. [17] provided the measurements for a number of popular DNN models on mobile CPUs and GPUs (Nvidia TK1 and TX1). Ran et al.[23] further explored the accuracy-latency tradeoffs on mobile devices by measuring how reducing the dimensionality of the input size reduces the overall accuracy and latency. DNN models designed specifically for mobile devices, such as MobileNets [14], report system performance in terms of a number of multiply-add operations, which could be used to estimate latency characteristics and other metrics on different mobile hardware, based on the processing capabilities of the hardware. Once the system performance is understood, the application developer can choose the right model.

2.1.4 Frameworks

Several open-source software libraries are publicly available for deep learning inference and training on end devices and edge servers. Google’s TensorFlow [4], released in 2015, is an interface for expressing machine learning algorithms and an implementation for executing such algorithms on heterogeneous distributed systems. Tensorflow’s computation workflow is designed as a directed graph and utilizes a placement algorithm to distribute computation tasks based on the estimated or measured execution time and communication time [5]. The placement algorithm uses a greedy

approach that places a computation task on the node that is expected to complete the computation the soonest. Tensorflow can run on edge devices, such as Raspberry Pi and smartphones. TensorFlow Lite was proposed in the late 2017 [3], which is an optimized version of Tensorflow for mobile and embedded devices, with mobile GPU support added in early 2019. Tensorflow Lite only provides on-device inference abilities, not training, and achieves low latency by compressing a pre-trained DNN model. Caffe [16] is another deep learning framework, originally developed by Jia, with the current version, Caffe2, maintained by Facebook. It seeks to provide an easy and straightforward way for deep learning with a focus on mobile devices, including smartphones and Raspberry Pis. PyTorch [20] is another deep learning platform developed by Facebook, with its main goal differing from Caffe2 in which it focuses on the integration of research prototypes to production development. Actually Facebook is working on the merge of Caffe2 and PyTorch frameworks. GPUs are an important key element in efficient DNN inference and training. NVIDIA provides GPU software libraries to make use of NVIDIA GPUs, such as CUDA [7] for general GPU processing and cuDNN [8] which is targeted toward deep learning. While such libraries are useful for training DNN models on a desktop server, cuDNN and CUDA are not widely available on current mobile devices such as smartphones. To utilize smartphone GPUs, Android developers can currently make use of Tensorflow Lite, which provides experimental GPU capabilities. To experiment with edge devices other than smartphones, researchers can turn to edge-specific development kits, such as the NVIDIA Jetson TX2 development kit for experimenting with edge computing, with NVIDIA-provided SDKs used to program the devices.

2.1.5 Challenges

Because of the required competences and effort to choose the best neural network and the best parameters that better fit the applications of our interest, there has also been much recent studies in automated machine learning, which uses artificial intelligence to choose which DNN model to run and tune the hyperparameters. For example, Tan et al. [29] and Taylor et al. [30] proposed using reinforcement learning and traditional machine learning, respectively, to choose the right hyperparameters for mobile devices, which is useful in edge scenarios. As described in [6] many challenges remain in

deploying deep learning on the edge, not only on end devices but also on the edge servers and on a combination of end devices, edge servers, and the cloud. For example parameters like latency, energy consumption and migration still the main challenges in the field of deep learning applied to the edge computing.

2.2 Edge Computing

Today, an IoT solution has to cover a much broader scope of requirements. We see that in most cases, organizations opt for a combination of cloud and edge computing for complex IoT solutions. Cloud computing typically comes into play when organizations require storage and computing power to execute certain applications and processes, and to visualize telemetry data from anywhere. Edge computing, on the other hand, is the right choice in cases with low latency, local autonomous actions, reduced back-end traffic, and when confidential data is involved.

2.2.1 Motivations

As written in [32] there are at least three reasons to consider the adoption of the edge computing. The first is that putting all the computing tasks on the cloud has been proved to be an efficient way for data processing since the computing power on the cloud outclasses the capability of the things at the edge. However, compared to the fast developing data processing speed, the bandwidth of the network has come to a standstill. With the growing quantity of data generated at the edge, speed of data transportation is becoming the bottleneck for the cloud-based computing paradigm. The second reason is that almost all kinds of electrical devices will become part of IoT, and they will play the role of data producers as well as consumers, such as air quality sensors, LED bars, streetlights and even an Internet-connected microwave oven. It is safe to infer that the number of things at the edge of the network will develop to more than billions in a few years. Thus, raw data produced by them will be enormous, making conventional cloud computing not efficient enough to handle all these data. This means most of the data produced by IoT will never be transmitted to the cloud, instead it will be consumed at the edge of the network. Fig. 2.2 shows the conventional cloud computing structure.




Figure 2.2: Cloud computing paradigm

However, this structure is not sufficient for IoT. First, data quantity at the edge is too large, which will lead to huge unnecessary bandwidth and computing resource usage. Second, the privacy protection requirement will pose an obstacle for cloud computing in IoT. Lastly, most of the end nodes in IoT are energy constrained things, and the wireless communication module is usually very energy hungry, so offloading some computing tasks to the edge could be more energy efficient. Another valid consideration to take into account is that in the cloud computing paradigm, the end devices at the edge usually play as data consumer, for example, watching a YouTube video on your smart phone. However, people are also producing data nowadays from their mobile devices. The change from data consumer to data producer/consumer requires more function placement at the edge.

2.2.2 Definitions

Edge computing refers to the enabling technologies allowing computation to be performed at the edge of the network, on downstream data on behalf of cloud services and upstream data on behalf of IoT services. The rationale of edge computing is that computing should happen at the proximity of data sources. From a certain point of view, edge computing could interchangeable with fog computing [13], but edge computing focus more toward the things side, while the first focus more on the infrastructure side. Fig. 2.3 illustrates the two-way computing streams in edge computing. In the edge computing paradigm, the things not only are data consumers, but also play as data producers. At the edge, the things can not only request service and content from the cloud but also perform the computing tasks from the cloud. Edge can perform computing offloading, data storage, caching and processing, as well as distribute request and delivery service from cloud to user. With those jobs in the network, the edge itself needs to be well designed to meet the requirement efficiently in service such as reliability, security, and privacy




Figure 2.3: Edge computing paradigm

protection.

2.2.3 Performance Metrics

In edge computing, we have multiple layers with different computation capability. Workload allocation becomes a big issue. We need to decide which layer to handle the workload or how many tasks to assign at each part. There are multiple allocation strategies to complete a workload, for instances, evenly distribute the workload on each layer or complete as much as possible on each layer. To choose an optimal allocation strategy, there are several optimization parameters to consider. These metrics are very similar to those described in 2.1, including latency, bandwidth and energy. How to measure these performance is strongly dependent by the environment, resources and application case.

2.2.4 Challenges

In cloud computing, users program their code and deploy them on the cloud. Usually, the program is written in one programming language and compiled for a certain target platform, since the program only runs in the cloud. However, in the edge computing, computation is offloaded from the cloud, and the edge nodes are most likely heterogeneous platforms. In this case, the runtime of these nodes differ from each other, and the programmer faces huge difficulties to write an application that may be deployed in the edge computing paradigm. Another important challenge is related to the *naming*. In edge computing, one important assumption is that the number of things is tremendously large. At the top of the edge nodes, there are a lot of applications running, and each application has its own structure about how the service is provided. Similar to all computer systems, the naming scheme in edge computing is very important for development,

addressing, things identification, and data communication. However, an efficient naming mechanism for the edge computing paradigm has not been built and standardized yet. Edge practitioners usually needs to learn various communication and network protocols in order to communicate with the heterogeneous things in their system. In terms of *service management* at the edge of the network, there are four fundamental features that should be supported to guarantee a reliable system, including differentiation, extensibility, isolation, and reliability. To protect the data *security* and usage *privacy* at the edge of the network, several challenges remain open. First is the awareness of privacy and security to the community: ip camera, health monitor, or even some WiFi enabled toys could easily be connected by others if not protected properly. Second is the ownership of the data collected from things at edge. Just as what happened with mobile applications, the data of end user collected by things will be stored and analyzed at the service provider side. Third is the missing of efficient tools to protect data privacy and security at the edge of the network. The highly dynamic environment at the edge of the network also makes the network become vulnerable or unprotected and more tools are still missing to handle diverse data attributes for edge computing.

2.3 Robotics Tools and Platforms

This section will cite and recap the main robotics platforms discovered and explored during the development of this thesis.

2.3.1 ROS

The Robot Operating System (ROS) [22], is a open-source framework based on the component-based software engineering paradigm that provides the middleware for inter-process communication. Initially developed by the Stanford Artificial Intelligence Laboratory its development continued at Willow Garage, a robotics research institute, and now it's maintained and improved under the action of the ORSF foundation [22]. As a meta-operating system, ROS offers features such as hardware abstraction, low-level device control, implementation of commonly used functionalities, message communication between processes and package management. It uses an asynchronous publish/subscribing mechanism made possible by message

standardisation and encapsulation that make the external interface of every node as general as possible, allowing quick nodes exchange and, thus, great architectural flexibility. Each independent block, called *node*, executes a particular task of a process and can communicate with other nodes through *topics*. This allows to create complex architectures by aggregating many simpler entities and simplifies the use of different tasks or different methods for the same task. Additionally to the message-passing system, the core ROS component, called *roscore*, maintains a global execution time for the nodes to achieve synchronisation. Each node executes separately with its own internal clock driven by the set execution rate. At every message sent/received, containing the internal time information, the core component updates the global time by following the execution status provided by these messages. For more clarity about how ROS communication works an example is provided in figure 2.4. When a publishing node will announce to the master that it is publishing over a topic and the subscribing node will say to the master that it want to listen to a topic. Then if the publisher and the subscriber are on the same topic, the master will transfer data in order for the two other node to communicate directly via TCP/IP or UDP.

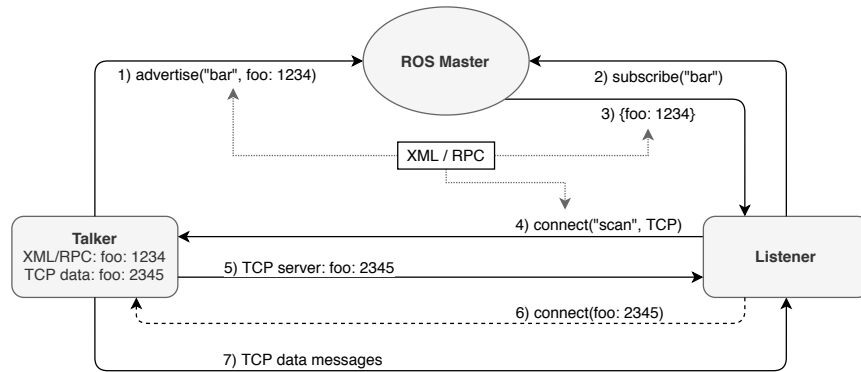


Figure 2.4: Example of nodes and topic communication

The reason to choose ROS for the development of robotics applications resides mainly in its high modularity. Furthermore it is a well known standard in the scientific community of the robotics researches and developers.

2.3.2 NVIDIA Isaac SDK

The Isaac SDK is the main software toolkit provided by NVIDIA and allows developers to brought new opportunities for developing robotics solutions as well as researching topics in this field. As described in [9] it's comprised of the following:

- Isaac Robot Engine: A framework which allows you to easily write modular applications and deploy them on your robots.
- Isaac GEMs: A collection of robotics algorithms from planning to perception, most of them GPU-accelerated.
- Applications: Various example applications from basic samples which show specific features to applications that facilitate complicated robotics use cases.
- Isaac Sim for Navigation: A powerful virtual robotics laboratory and a high-fidelity 3D world simulator that accelerates research, design, and development by reducing cost and risk.

The main reasons we would consider to use this platform reside in its high modularity and high performance solutions that it is able to offer on NVIDIA Jetson platforms. Others important reasons to adopt this solution are related the **C API** that allows the communication with Isaac apps from languages other than C++, and the **ROS bridge** that offers the capability to interact with ROS nodes applications.

2.3.3 CoppeliaSim

From the creators of V-Rep [10], CoppeliaSim [25] seems to be the definitive solution for development of robotics applications. The robot simulator CoppeliaSim, with integrated development environment, is based on a distributed control architecture: each object/model can be individually controlled via an embedded script, a plugin, ROS nodes, BlueZero nodes, remote API clients, or a custom solution. This makes CoppeliaSim very versatile and ideal for multi-robot applications. Controllers can be written in C/C++, Python, Java, Lua, Matlab, Octave or Urbi. CoppeliaSim can be used as a stand-alone application or can easily be embedded into a main client application: its small footprint and elaborate API makes CoppeliaSim

an ideal candidate to embed into higher-level applications. An integrated Lua script interpreter makes CoppeliaSim an extremely versatile application, leaving the freedom to the user to combine the low/high-level functionalities to obtain new high-level functionalities.

2.4 Docker

While container technologies have been around for a long time, they've become more widely known with the rise of the Docker container platform. Docker was the first container system that made containers easily portable across different machines. It simplified the process of packaging up not only the application but also all its libraries and other dependencies, even the whole OS file system, into a simple, portable package that can be used to provision the application to any other machine running Docker. When you run an application packaged with Docker, it sees the exact filesystem contents that you've bundled with it. It sees the same files whether it's running on your development machine or a production machine, even if the production server is running a completely different Linux OS. The application won't see anything from the server it's running on, so it doesn't matter if the server has a completely different set of installed libraries compared to your development machine. This is similar to creating a VM image by installing an operating system into a VM, installing the app inside it, and then distributing the whole VM image around and running it. Docker achieves the same effect, but instead of using VMs to achieve app isolation, it uses Linux container technologies to provide (almost) the same level of isolation that VMs do. Instead of using big monolithic VM images, it uses container images, which are usually smaller.

2.4.1 Understanding Docker Concepts

Docker is a platform for packaging, distributing, and running applications. As we've already stated, it allows you to package your application together with its whole environment. This can be either a few libraries that the app requires or even all the files that are usually available on the filesystem of an installed operating system. Docker makes it possible to transfer this package to a central repository from which it can then be transferred to any computer running Docker and executed there. Three main concepts

in Docker comprise this scenario:

- **Images:** A Docker-based container image is something you package your application and its environment into. It contains the filesystem that will be available to the application and other metadata, such as the path to the executable that should be executed when the image is run.
- **Registries:** A Docker Registry is a repository that stores your Docker images and facilitates easy sharing of those images between different people and computers. When you build your image, you can either run it on the computer you've built it on, or you can push (upload) the image to a registry and then pull (download) it on another computer and run it there. Certain registries are public, allowing anyone to pull images from it, while others are private, only accessible to certain people or machines.
- **Containers:** A Docker-based container is a regular Linux container created from a Docker-based container image. A running container is a process running on the host running Docker, but it's completely isolated from both the host and all other processes running on it. The process is also resource-constrained, meaning it can only access and use the amount of resources (CPU, RAM, and so on) that are allocated to it.

Figure 2.5 shows all three concepts and how they relate to each other. The developer first builds an image and then pushes it to a registry. The image is thus available to anyone who can access the registry. They can then pull the image to any other machine running Docker and run the image. Docker creates an isolated container based on the image and runs the binary executable specified as part of the image. Docker images are composed of layers. Different images can contain the exact same layers because every Docker image is built on top of another image and two different images can both use the same parent image as their base. This speeds up the distribution of images across the network, because layers that have already been transferred as part of the first image don't need to be transferred again when transferring the other image. But layers don't only make distribution more efficient, they also help reduce the storage footprint of images. Each

layer is only stored once. Two containers created from two images based on the same base layers can therefore read the same files, but if one of them writes over those files, the other one doesn't see those changes. Therefore, even if they share files, they're still isolated from each other. This works because container image layers are read-only. When a container is run, a new writable layer is created on top of the layers in the image. When the process in the container writes to a file located in one of the underlying layers, a copy of the whole file is created in the top-most layer and the process writes to the copy.

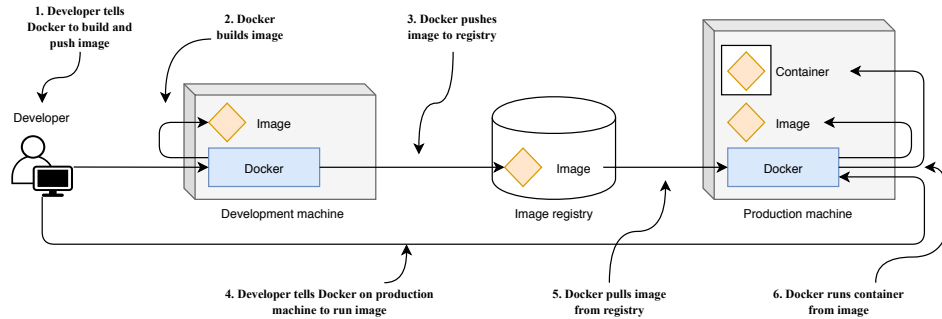


Figure 2.5: Docker images, registries, and containers

2.4.2 Portability Limitations of Container Images

In theory, a container image can be run on any Linux machine running Docker, but one small caveat exists, one related to the fact that all containers running on a host use the host's Linux kernel. If a containerized application requires a specific kernel version, it may not work on every machine. If a machine runs a different version of the Linux kernel or doesn't have the same kernel modules available, the app can't run on it.

While containers are much more lightweight compared to VMs, they impose certain constraints on the apps running inside them. VMs have no such constraints, because each VM runs its own kernel.

And it's not only about the kernel. It should also be clear that a containerized app built for a specific hardware architecture can only run on other machines that have the same architecture. It isn't possible to containerize an application built for the x86 architecture and expect it to run on an ARM-based machine because it also runs Docker. You still need a VM for

that. It is very important understand that Docker itself doesn't provide process isolation. The actual isolation of containers is done at the Linux kernel level using kernel features such as Linux Namespaces and cgroups. Docker only makes it easy to use those features.

2.5 Kubernetes

As the number of deployable application components in your system grows, it becomes harder to manage them all. Google was probably the first company that realized it needed a much better way of deploying and managing their software components and their infrastructure to scale globally. It's one of only a few companies in the world that runs hundreds of thousands of servers and has had to deal with managing deployments on such a massive scale. This has forced them to develop solutions for making the development and deployment of thousands of software components manageable and cost-efficient.

2.5.1 What is Kubernetes

Kubernetes is a software system that allows you to easily deploy and manage containerized applications on top of it. It relies on the features of Linux containers to run heterogeneous applications without having to know any internal details of these applications and without having to manually deploy these applications on each host. Because these apps run in containers, they don't affect other apps running on the same server, which is critical when you run applications for completely different organizations on the same hardware. Kubernetes enables you to run your software applications on thousands of computer nodes as if all those nodes were a single, enormous computer. Deploying applications through Kubernetes is always the same, whether your cluster contains only a couple of nodes or thousands of them. The size of the cluster makes no difference at all. Additional cluster nodes simply represent an additional amount of resources available to deployed apps.

Figure 2.6 shows the simplest possible view of a Kubernetes system. The system is composed of a master node and any number of worker nodes. When the developer submits a list of apps to the master, Kubernetes deploys them to the cluster of worker nodes. What node a component

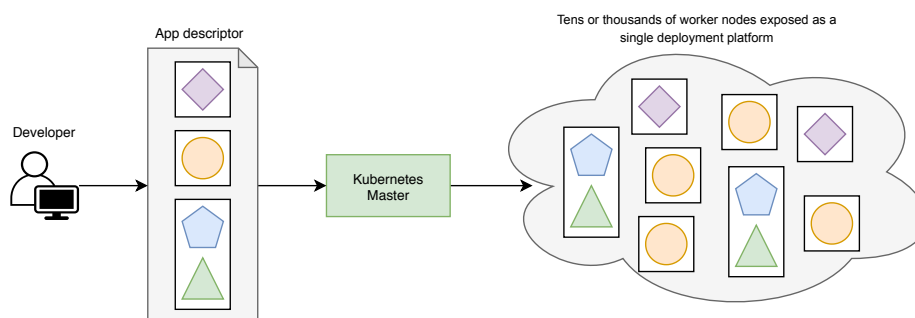


Figure 2.6: Kubernetes exposes the whole datacenter as a single deployment platform

lands on doesn't (and shouldn't) matter—neither to the developer nor to the system administrator.

2.5.2 Architecture of a Kubernetes cluster

At the hardware level, a Kubernetes cluster is composed of many nodes, which can be split into two types. The first is the *Control Plane* that is what controls the cluster and makes it function. It consists of multiple components that can run on a single master node or be split across multiple nodes and replicated to ensure high availability. These components are:

- The *Kubernetes API Server*, which you and the other Control Plane components communicate with.
- The *Scheduler*, which schedules your apps (assigns a worker node to each deployable component of your application).
- The *Controller Manager*, which performs cluster-level functions, such as replicating components, keeping track of worker nodes, handling node failures, and so on.
- *etcd*, a reliable distributed data store that persistently stores the cluster configuration.

The components of the Control Plane hold and control the state of the cluster, but they don't run your applications. This is done by the (worker) nodes. The worker nodes are the machines that run your containerized

applications. The task of running, monitoring, and providing services to your applications is done by the following components:

- Docker, rkt, or another container runtime, which runs your containers.
- The Kubelet, which talks to the API server and manages containers on its node.
- The Kubernetes Service Proxy (kube-proxy), which load-balances network traffic between application components.

Figure 2.7 shows the components running on these two sets of nodes.

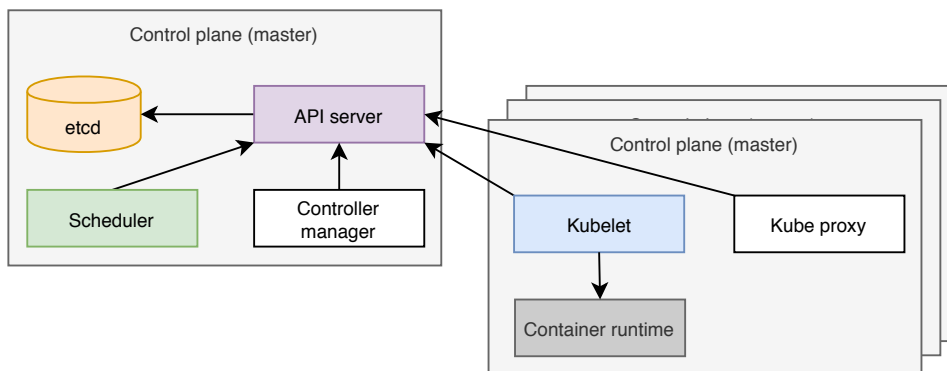


Figure 2.7: The components that make up a Kubernetes cluster

2.5.3 Running an application in Kubernetes

To run an application in Kubernetes, you first need to package it up into one or more container images, push those images to an image registry, and then post a description of your app to the Kubernetes API server.

The description includes information such as the container image or images that contain your application components, how those components are related to each other, and which ones need to be run co-located (together on the same node) and which don't. For each component, you can also specify how many copies (or replicas) you want to run. Additionally, the description also includes which of those components provide a service to either internal or external clients and should be exposed through a single IP address and made discoverable to the other components. When the API server processes your app's description, the Scheduler schedules the specified groups of containers onto the available worker nodes based on computational resources

required by each group and the unallocated resources on each node at that moment. The Kubelet on those nodes then instructs the Container Runtime (Docker, for example) to pull the required container images and run the containers.

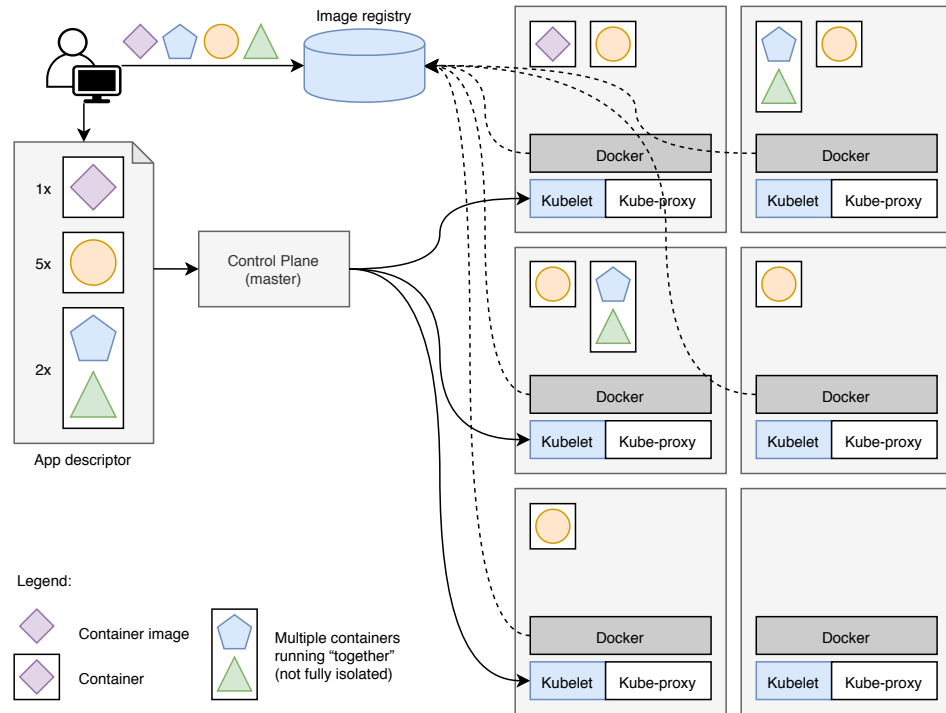


Figure 2.8: A basic overview of the Kubernetes architecture and an application running on top of it.

Examine figure 2.8 to gain a better understanding of how applications are deployed in Kubernetes. The app descriptor lists four containers, grouped into three sets (these sets are called pods). The first two pods each contain only a single container, whereas the last one contains two. That means both containers need to run co-located and shouldn't be isolated from each other. Next to each pod, you also see a number representing the number of replicas of each pod that need to run in parallel. After submitting the descriptor to Kubernetes, it will schedule the specified number of replicas of each pod to the available worker nodes. The Kubelets on the nodes will then tell Docker to pull the container images from the image registry and run the containers.

2.6 Kubeedge

Chapter 3

Verification

This chapter will focus on description and verification of ORB-SLAM2 [18] improved solution proposed in the IROS 2019 conference [28]. Primary it will describe how ORB-SLAM2 works. Secondly it will show an accurate inspection necessary to find some potential bugs that cause an unexpected behaviour of the algorithm in a sequential environment.

3.1 System description

As described in [28], ORB-SLAM2 is a Simultaneous Localization And Mapping system that can work with data coming from monocular, stereo, and RGB-D cameras. A system of this kind has the purpose to allow both map reconstruction and navigation in the most common environment without the support of a GPS. The system consists of the following three main blocks (see figure 3.1):

Tracking and localization This block is in charge of computing visual features, localizing the robot in the environment, and, in case of significant discrepancies between an already saved map and the input stream, communicating updated map information to the mapping block. The frames per second (FPS) that can be computed by the whole system strongly depends on the performance of this block.

Mapping It updates the environment map by using the information (map changes) sent by the localization block. It is a computational time consuming block and its execution rate strictly depends on the agent speed. However, considering the actual agent speed of the KITTI

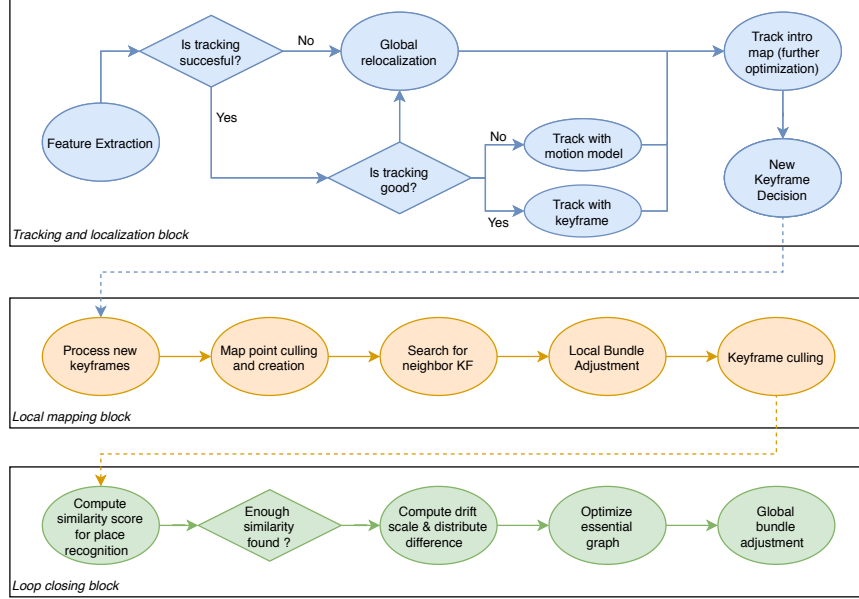


Figure 3.1: Main blocks of the ORB-SLAM2 algorithm.

datasets analysed in this work [11], it does not represent a system bottleneck.

Loop closing It aims at adjusting the scale drift error accumulated during the input analysis. When a loop in the robot pathway is detected, this block updates the mapped information through a high latency heavy computation, during which the first two blocks are suspended. This can lead the robot to loose tracking and localization information and, as a consequence, the robot to get temporary lost. The computation efficiency of this block (running on-demand) is crucial for the quality of the final results.

The system is organized on three parallel threads, one per block. The use of parallel threads allows for obtaining real-time processing on an Intel Core i7 desktop PC with 16GB RAM [18]. The original open source version of ORB-SLAM2 provides two level of parallelism. The first level is given by the three main algorithm blocks (see Fig. 3.1), which are implemented to be run as parallel PThreads on shared-memory multi-core CPUs. The second level is given by the automatic parallel implementation (i.e., through OpenMP directives) of the *bundle adjustment* sub-block, which is part both of the local mapping and loop closing blocks. This allows the parallel com-




Figure 3.2: DAG of the feature extraction block and the corresponding sub-block implementations (GPU vs. CPU).

putation of such a long latency task on multi-core CPUs. To fully exploit the heterogeneous nature of the target NVIDIA Jetson TX2 board, the work done in [28] added two further levels of parallelism. The first is given by the parallel implementation for GPU of a set of tracking sub-blocks (see Fig. 3.2). This is because the most important bottleneck that characterizes the processing rate in terms of supported FPS, was detected on the feature extraction block. The second is given by the implementation of a 8-stage pipeline of such sub-blocks in order to benefit of an overlapped computation.

3.2 Unexpected behaviour

After the porting and the optimizations of the ORB-SLAM2 algorithm on the NVIDIA Jetson TX2 to achieve real-time performance, the experimental results conducted in [28] shown a massive improvement of the execution times on some KITTI dataset sequences. Despite these performance improvements, the algorithm seems to exhibit some unexpected behaviours too frequently, especially with sequences that contain a higher space-temporal variation between one frame and the next one. More specifically, it get into the Tack Lost state (fig. 3.3a) in a totally randomic way. Another Track lost state is reached when less then 5 Keyframes are computed and it loses its position (see fig: 3.3b).

Due to the randomness with which the problem occurs, to identify the cause it's necessary isolate the problem as much as possible. For this purpose



(a) Track lost state



(b) Track lost soon after initialization, system resetting.

Figure 3.3: LOST states of ORB-SLAM2

the first step has been sequentialize the execution between Tracking and Mapping blocks.

3.2.1 The problem of the non-determinism

3.2.2 The problem of the concurrency

3.3 Discussion

In this chapter the problem of non deterministic behaviour has been addressed and some of the possible solutions to overcome the problem of ... have been presented.

The strategy adopted to achieve the goal of ... could reduce the ... in the system.

Chapter 4

Methodology

This chapter will present a multi-level model design flow for the integration of heterogeneous applications and their deployment to the edge on embedded low power devices like NVIDIA Jetson TX2.

4.1 A ROS based robotic system

The choice of the most suitable integration platform is one of the main problems to be faced when integrating heterogeneous applications such as robotic software systems. This choice will have implications on performance and communication of the whole system. Furthermore it will be a key factor during the design phase of the entire architecture. After an accurate comparison among all available robotics platforms, for the purpose of the specific application developed in this thesis it has been decided to use ROS [22] in favour of its standard on communication among other robotic applications.

The design flow proposed in this work is composed of three bottom-up architecture levels, but the same methodology can be applied with more levels. The purpose of this methodology is to bring an heterogeneous application within an orchestration system such as Kubernetes, and then extend it to the edge computing. The architecture are explained in the next sections.

4.2 First level architecture

This is the first architecture to start with. At this architecture level the purpose is to verify the functioning of the whole system with the only

ROS capabilities on a single powerful machine. A simple scheme of this architecture is shown in figure 4.1.

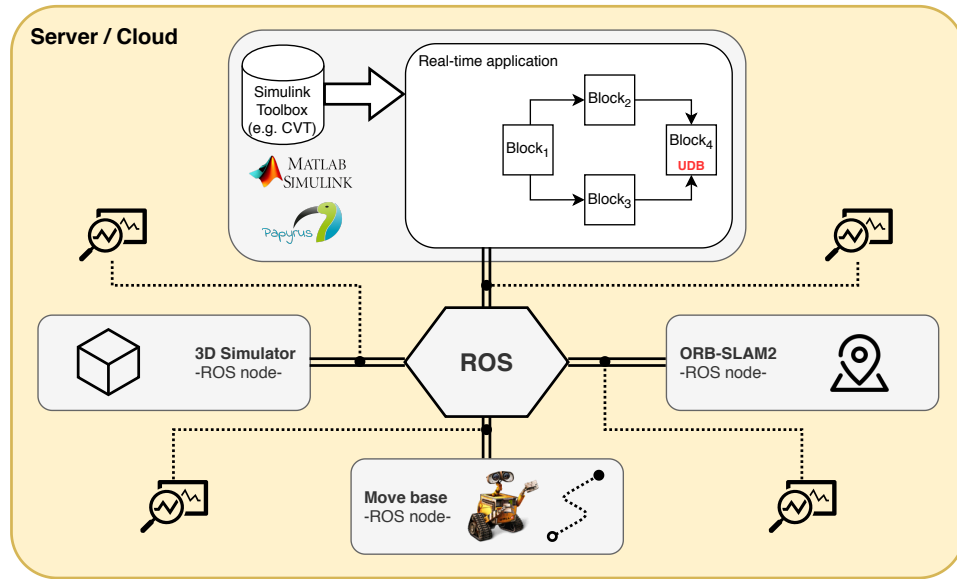


Figure 4.1: L1 architecture

The first step to realize this architecture is to wrap your (heterogeneous) applications with the ROS conformance protocol *advertises* and *callbacks*. An example of this operation is provided in listings 4.1 and 4.2. In this way, all the applications that compose your system will be able to communicate among them following the standard communication based on ROS. The ROS nodes (publishers and subscribers) communicate between them through topics, thus it is necessary choice the right names of these topics and who will read/write on them. Once the applications are ready to run on ROS, they can be executed using a launchfile, where the user can specify a set of nodes to launch and their parameters. Another possibility is invoke the `roslaunch` command for each node that have to be executed.

Listing 4.1: Publisher

```

#include "ros/ros.h"
...
int main(int argc, char **argv) {
    ros::init(argc, argv, "nodeA");
    ros::NodeHandle n;
    ros::Publisher pub;
    pub=n.advertise<T>("topicA",10);
    while (ros::ok()) {
        <put here your application>
        ...
        pub.publish(msg);
    }
    return 0;
}

```

Listing 4.2: Subscriber

```

#include "ros/ros.h"
...
void subCallback(const <T> msg) {
    <do stuff with received message>
}

int main(int argc, char **argv) {
    ros::init(argc, argv, "nodeB");
    ros::NodeHandle n;
    ros::Subscriber sub;
    sub=n.subscribe("topicA",10,subCall);
    ...
    return 0;
}

```

Commands such as `rostopic nodes` and `rostopic list` can be used to verify the correct communication between the nodes that compose the system. There exist more other tools to better inspect an application based on ROS, and they are all available online on ROS documentation. A good practise for performance measurement in ROS is to use profilers such as: `perf`, `gprof` and `callgrind`. `Rqt_top` is another profiler tool available in ROS, but it doesn't seem very stable and precise yet. Nevertheless, the best choice still put some time points in your code and write the measured intervals on files.

4.3 Second level architecture

Once the system is working properly on a single x86 host machine, it is possible to split it deploying manually the ROS nodes on a low power embedded device. Again, it is necessary pay attention on communication between nodes, because now they are on different machines and they need to communicate between them.

It would be preferable to have the control over which nodes run on which machines, for load-balancing and bandwidth management. Intuitively a user would duplicate the `roslaunch` files and comment them in a mutually exclusive way, such that the nodes that run on a machine, don't run on the other. Even if it could work correctly, this kind of hard-coding is not the best

practice for this purpose, especially for large projects. The recommended way for this purpose, is taking advantage of the tag machine explained in the ROS documentation.

If the roslaunch files have been properly defined, the only thing to do is ...

4.4 Third level architecture

Chapter 5

Experimental Results

Now it will be presented the experimental results obtained using the methodology explained in the previous chapter. First it will be described the setup Hardware and Software used during the experiments. Next, it will be explained how the methodology can be applied considering a typical robotic system developed for mobile robot navigation in a warehouse.

5.1 Setup

The main technicals specifications of the *Host* are shown in table 5.1.

Hardware	Model
Processor	Intel(R) Core(TM) i5-7400 CPU @ 3.00GHz
Operating System	18.04.1-Ubuntu
Storage	SATA 500GiB
Memory	16GiB
GPU	NVIDIA GeForce GTX 980
Software	Version
OpenCV	3.4.7
CUDA	10.0
ROS	Melodic Moreina

Table 5.1: Technical specifications of the Host

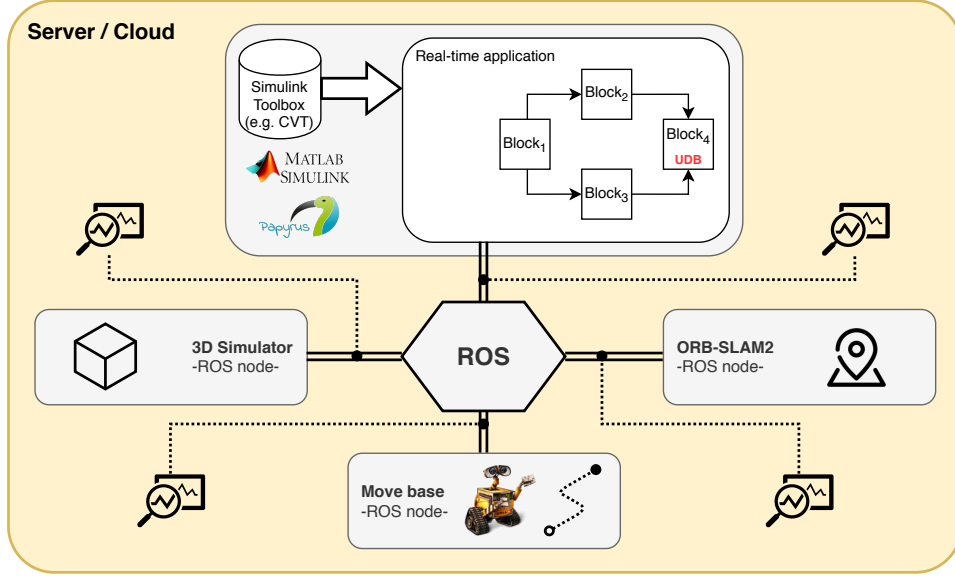


Figure 5.1: L1 architecture

5.2 Mobile robots

5.2.1 Architecture at L1

At this architecture level we put a whole ROS-based system on a single Intel x86 powered machine that we will conventionally call *Host*. This is the simplest and the most common architecture used by robotics researches. As shown in figure 5.1, in this case there are two main ROS nodes that communicate between them. Specifically speaking, there are both ORB-SLAM2, explained in the previous chapter, and a hybrid motion planner (VOVD [21]). Both applications are running as ROS nodes to perform a mobile robot 2D navigation task.

Because of the most used robotics platforms used today support the Melodic Morenia [2] version of ROS, and because of VOVd was written with the Kinetic version, it has been decided to convert VOVd in the relative Melodic ROS version. This task required a redefinition of some functions related to the ROS *tf2* package [31] due to the deprecation of the *tf* package functions previously used. Another modification done talking about VOVd is related to the 3D map used by Gazebo [1]. The problem was that it came with a featureless map generated from an URDF file, and because of ORB-SLAM2 properly works only if some features are detected during

its execution, it has been necessary to add a material to the 3D model of the map. Talking about ORB-SLAM2, it comes already provided with the necessary functions to run in a ROS environment, thus no other operations have been required to complete our first architecture level.

5.2.2 Architecture at L2

Another common architecture in robotics, involves the distribution of ROS nodes on different hardware platforms (e.g. NVIDIA Jetson or Raspberry). This step requires the split of our system in two parts: *Host* (with Gazebo simulator) and *Device/s* (robot navigation node/s). For this purpose, due to the several interconnected nodes, each of which having many parameters, it is recommended to create a roslaunch file using the *machine tag*. These tags allow control over which nodes run on which machines, for load-balancing and bandwidth management. Fig. 5.2 shows an example of this configuration. At this architecture level the main focus is on the communication between the *Host*, that act as Master, and the other external nodes that represent the slaves. It is important to verify that the system continues working properly and that the Network doesn't become a bottleneck. For example, if we need to transfer a stream of images from one ROS node to another that is running on a different machine, we should compress the images before sending them over the network. In this way we avoid saturating the bandwidth of the network.

5.2.3 Architecture at L3

Finally the last typical architecture level is composed by *Host*, *Device/s* and *Robot/s*. If the previous steps have been properly done, this one should be completely transparent with respect to the entire system. As we can see in Fig. 5.3, the simulator has been removed from the *Host* and it is replaced by the real robot, on which one or more embedded devices are attached. To achieve this goal ...

5.3 Deployment from the cloud to the edge

As mentioned in 2.6, Kubeedge is built upon Kubernetes and provides fundamental infrastructure support for network, app. deployment and meta-data synchronization between cloud and edge. As we know, Kubernetes uses

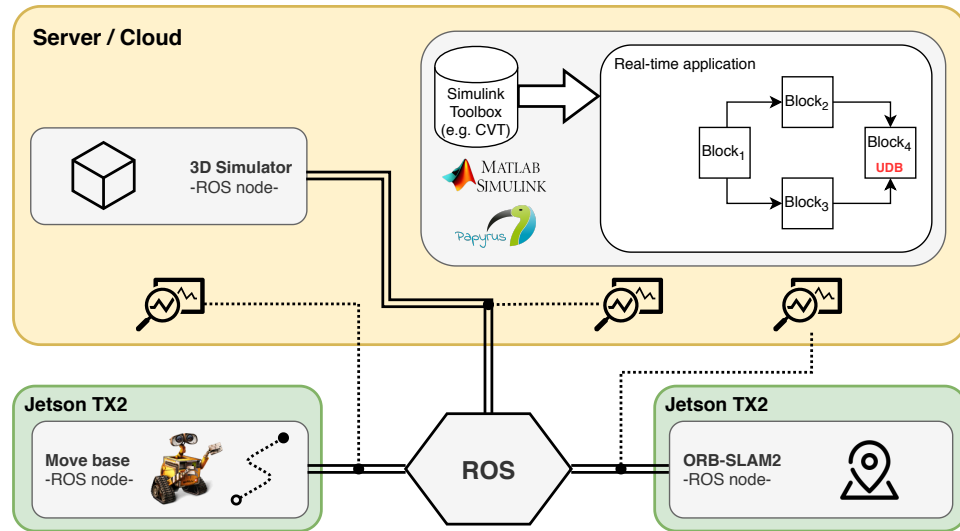


Figure 5.2: L2 architecture

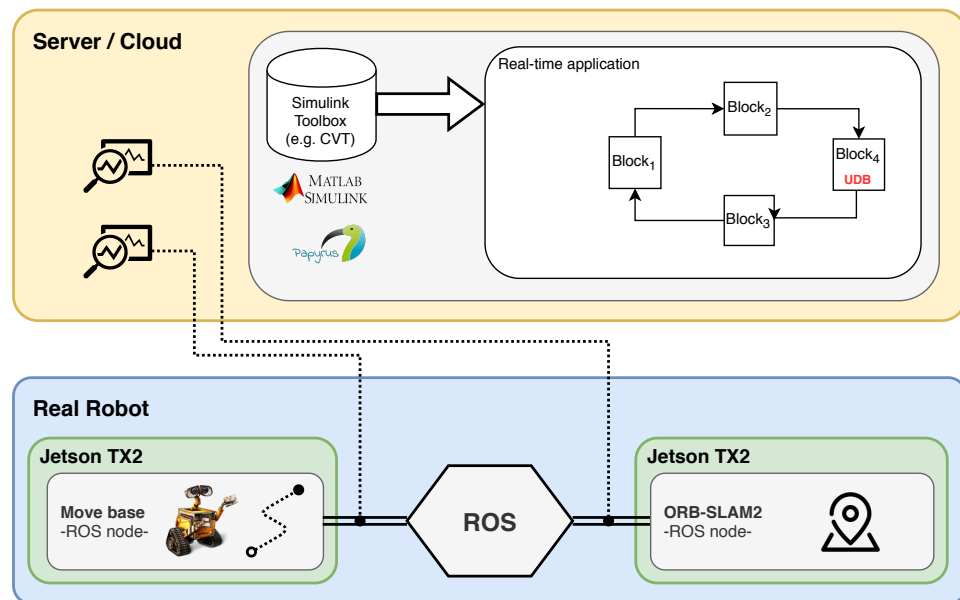


Figure 5.3: L3 architecture

containers to run isolated, packaged applications across its cluster nodes. To run on Kubernetes, your applications must be encapsulated in one or more container images and executed using a container runtime like Docker. While containerizing your components is a requirement for Kubernetes, it also allows easy scaling and management. For instance, containers provide isolation between the application environment and the external host system, support a networked, service-oriented approach to inter-application communication, and typically take configuration through environmental variables and expose logs written to standard error and standard out. Containers themselves encourage process-based concurrency and help maintain dev/prod parity by being independently scalable and bundling the process's runtime environment. These characteristics make it possible to package your applications so that they run smoothly on Kubernetes. In this section it will be described the process used in our project to prepare both VOVD and ORB-SLAM2 applications for the edge computing through the support of Kubeedge.

5.3.1 VOVD at the edge

5.3.2 ORB-SLAM2 at the edge

5.4 The whole system on Kubeedge

5.5 Discussion

5.5.1 Hybrid local/global planner

5.6 Results

Setup	Odom VOVD (ms)	Track Pipe (m)
orb(j1)	-	43.41
vovd(j1)	1.70	-
(orb;vovd)(j1)	2.58	73.49
(orb(j1);vovd(j2))	1.60	49.81
Performance Measurement: camera 20 Hz		
orb(j1)	-	54.82
vovd(j1)	1.70	-
(orb;vovd)(j1)	2.85	93.15
(orb(j1);vovd(j2))	1.60	56.20
Kube		
orb(j1)	-	47.28
vovd(j1)	1.30	-
(orb;vovd)(j1)	1.40	82.47
(orb(j1);vovd(j2))	1.38	50.83
Kube		
orb(j1)	-	50.20
vovd(j1)	1.21	-
(orb;vovd)(j1)	1.43	70.54
(orb(j1);vovd(j2))	1.26	58.14

Table 5.2: Performance Measurement

Setup	Odom (ms)	Track Pipe (ms)	Dec. (ms)	Frame Elab. (ms)	Feature Ext.(ms)	Supp. (FPS)
Performance Measurement: camera 10 Hz						
orb(j1)	-	43.41	13.08	49.21	43.62	20.32
vovd(j1)	1.70	-	-	-	-	-
(orb;vovd)(j1)	2.58	73.49	18.69	62.70	54.76	13.61
(orb(j1);vovd(j2))	1.60	49.81	13.53	49.25	43.01	20.08
Performance Measurement: camera 20 Hz						
orb(j1)	-	54.82	14.47	49.40	44.35	18.24
vovd(j1)	1.70	-	-	-	-	-
(orb;vovd)(j1)	2.85	93.15	23.96	67.99	60.41	10.74
(orb(j1);vovd(j2))	1.60	56.20	14.00	49.32	43.95	17.79
Kubeedge - Performance Measurement: camera 10 Hz						
orb(j1)	-	47.28	12.78	40.52	36.02	21.15
vovd(j1)	1.30	-	-	-	-	-
(orb;vovd)(j1)	1.40	82.47	15.10	53.54	48.09	12.12
(orb(j1);vovd(j2))	1.38	50.83	13.98	43.80	38.04	19.67
Kubeedge - Performance Measurement: camera 20 Hz						
orb(j1)	-	50.20	16.20	42.16	34.82	19.92
vovd(j1)	1.21	-	-	-	-	-
(orb;vovd)(j1)	1.43	70.54	15.02	55.60	50.24	14.18
(orb(j1);vovd(j2))	1.26	58.14	13.52	42.34	37.17	17.20

Table 5.3: Performance Measurement.

Conclusions And Future Works

In this thesis a has been presented to address the issue of providing
...

Bibliography

- [1] Gazebo robot simulation. <http://gazebo-sim.org/>.
- [2] Ros melodic morenia. <http://wiki.ros.org/melodic>.
- [3] Tensorflow lite. <https://www.tensorflow.org/lite>.
- [4] Martín Abadi, Ashish Agarwal, Paul Barham, Eugene Brevdo, Zhifeng Chen, Craig Citro, Greg S. Corrado, Andy Davis, Jeffrey Dean, Matthieu Devin, Sanjay Ghemawat, Ian Goodfellow, Andrew Harp, Geoffrey Irving, Michael Isard, Yangqing Jia, Rafal Jozefowicz, Lukasz Kaiser, Manjunath Kudlur, Josh Levenberg, Dandelion Mané, Rajat Monga, Sherry Moore, Derek Murray, Chris Olah, Mike Schuster, Jonathon Shlens, Benoit Steiner, Ilya Sutskever, Kunal Talwar, Paul Tucker, Vincent Vanhoucke, Vijay Vasudevan, Fernanda Viégas, Oriol Vinyals, Pete Warden, Martin Wattenberg, Martin Wicke, Yuan Yu, and Xiaoqiang Zheng. TensorFlow: Large-scale machine learning on heterogeneous systems, 2015. Software available from tensorflow.org.
- [5] Martín Abadi, Ashish Agarwal, Paul Barham, Eugene Brevdo, Zhifeng Chen, Craig Citro, Greg S. Corrado, Andy Davis, Jeffrey Dean, Matthieu Devin, Sanjay Ghemawat, Ian Goodfellow, Andrew Harp, Geoffrey Irving, Michael Isard, Yangqing Jia, Rafal Jozefowicz, Lukasz Kaiser, Manjunath Kudlur, Josh Levenberg, Dan Mane, Rajat Monga, Sherry Moore, Derek Murray, Chris Olah, Mike Schuster, Jonathon Shlens, Benoit Steiner, Ilya Sutskever, Kunal Talwar, Paul Tucker, Vincent Vanhoucke, Vijay Vasudevan, Fernanda Viegas, Oriol Vinyals, Pete Warden, Martin Wattenberg, Martin Wicke, Yuan Yu, and Xiaoqiang Zheng. Tensorflow: Large-scale machine learning on heterogeneous distributed systems, 2016.

- [6] J. Chen and X. Ran. "deep learning with edge computing: A review". in *Proceedings of the IEEE*, 107(8):1655–1674, 2019.
- [7] NVIDIA Corporation. Cuda zone. <https://developer.nvidia.com/cuda-zone>.
- [8] NVIDIA Corporation. cuDNN deep neural network library. <https://developer.nvidia.com/cudnn>.
- [9] NVIDIA Corporation. Isaac sdk. <https://developer.nvidia.com/isaac-sd>.
- [10] Marc Freese, Surya Singh, Fumio Ozaki, and Nobuto Matsuhira. Virtual robot experimentation platform v-rep: A versatile 3d robot simulator. In *Proceedings of the Second International Conference on Simulation, Modeling, and Programming for Autonomous Robots*, SIMPAR'10, page 51–62, Berlin, Heidelberg, 2010. Springer-Verlag.
- [11] Andreas Geiger. Are we ready for autonomous driving? the kitti vision benchmark suite. In *Proceedings of the 2012 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, CVPR '12, page 3354–3361, USA, 2012. IEEE Computer Society.
- [12] Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep Learning*. MIT Press, 2016. <http://www.deeplearningbook.org>.
- [13] OpenFog Consortium Architecture Working Group. Openfog reference architecture for fog computing. https://www.iiconsortium.org/pdf/OpenFog_Reference_Architecture_2_09_17.pdf.
- [14] Andrew G. Howard, Menglong Zhu, Bo Chen, Dmitry Kalenichenko, Weijun Wang, Tobias Weyand, Marco Andreetto, and Hartwig Adam. Mobilenets: Efficient convolutional neural networks for mobile vision applications, 2017.
- [15] Jonathan Huang, Vivek Rathod, Chen Sun, Menglong Zhu, Anoop Korattikara, Alireza Fathi, Ian Fischer, Zbigniew Wojna, Yang Song, Sergio Guadarrama, and Kevin Murphy. Speed/accuracy trade-offs for modern convolutional object detectors, 2016.

- [16] Yangqing Jia, Evan Shelhamer, Jeff Donahue, Sergey Karayev, Jonathan Long, Ross Girshick, Sergio Guadarrama, and Trevor Darrell. Caffe: Convolutional architecture for fast feature embedding. In *Proceedings of the 22nd ACM International Conference on Multimedia*, MM '14, page 675–678, New York, NY, USA, 2014. Association for Computing Machinery.
- [17] Zongqing Lu, Swati Rallapalli, Kevin Chan, and Thomas La Porta. Modeling the resource requirements of convolutional neural networks on mobile devices. In *Proceedings of the 25th ACM International Conference on Multimedia*, MM '17, page 1663–1671, New York, NY, USA, 2017. Association for Computing Machinery.
- [18] Raúl Mur-Artal and Juan D. Tardós. ORB-SLAM2: an open-source SLAM system for monocular, stereo and RGB-D cameras. *IEEE Transactions on Robotics*, 33(5):1255–1262, 2017.
- [19] Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. Bleu: A method for automatic evaluation of machine translation. In *Proceedings of the 40th Annual Meeting on Association for Computational Linguistics*, ACL '02, page 311–318, USA, 2002. Association for Computational Linguistics.
- [20] Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, Alban Desmaison, Andreas Kopf, Edward Yang, Zachary DeVito, Martin Raison, Alykhan Tejani, Sasank Chilamkurthy, Benoit Steiner, Lu Fang, Junjie Bai, and Soumith Chintala. Pytorch: An imperative style, high-performance deep learning library. In H. Wallach et al., editor, *Advances in Neural Information Processing Systems 32*, pages 8024–8035. Curran Associates, Inc., 2019.
- [21] Nicola Piccinelli and Riccardo Muradore. Hybrid motion planner integrating global voronoi diagrams and local velocity obstacle method. pages 26–31, 06 2018.
- [22] Morgan Quigley, Ken Conley, Brian Gerkey, Josh Faust, Tully Foote, Jeremy Leibs, Rob Wheeler, and Andrew Ng. Ros: an open-source robot operating system. volume 3, 01 2009.

- [23] Xukan Ran, Haolanz Chen, Xiaodan Zhu, Zhenming Liu, and Jiasi Chen. Deepdecision: A mobile deep learning framework for edge video analytics. pages 1421–1429, 04 2018.
- [24] J. Redmon and A. Farhadi. Yolo9000: Better, faster, stronger. *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 6517–6525, 2017.
- [25] "E. Rohmer, S. P. N. Singh, and M. Freese". "coppeliasim (formerly v-rep): a versatile and scalable robot simulation framework". In "*Proc. of The International Conference on Intelligent Robots and Systems (IROS)*", "2013". "www.coppeliarobotics.com".
- [26] Sebastian Ruder. An overview of gradient descent optimization algorithms., 2016. cite arxiv:1609.04747Comment: Added derivations of AdaMax and Nadam.
- [27] Olga Russakovsky, Jia Deng, Hao Su, Jonathan Krause, Sanjeev Satheesh, Sean Ma, Zhiheng Huang, Andrej Karpathy, Aditya Khosla, Michael Bernstein, Alexander C. Berg, and Li Fei-Fei. ImageNet Large Scale Visual Recognition Challenge. *International Journal of Computer Vision (IJCV)*, 115(3):211–252, 2015.
- [28] D. D. Bloisi S. Aldegheri, N. Bombieri and A. Farinelli. Data flow orb-slam for real-time performance on embedded gpu boards. *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 5370–5375, 2019.
- [29] Mingxing Tan, Bo Chen, Ruoming Pang, Vijay Vasudevan, Mark Sandler, Andrew Howard, and Quoc V. Le. Mnasnet: Platform-aware neural architecture search for mobile, 2018.
- [30] Ben Taylor, Vicent Sanz Marco, Willy Wolff, Yehia Elkhatib, and Zheng Wang. Adaptive selection of deep learning models on embedded systems, 2018.
- [31] Wim Meeussen Tully Foote, Eitan Marder-Eppstein. tf ros package. <http://wiki.ros.org/tf>.
- [32] Q. Zhang Y. Li W. Shi, J. Cao and L. Xu. Edge computing: Vision and challenges. in *IEEE Internet of Things*, "3"("5"):637–646, 2016.